



Fastweb AI Agents & AI Factory

Fastweb provides **open-source AI models** hosted on **FastCloud infrastructure**, designed to cover most use cases in **Generative AI, Conversational AI, Multimodal AI, and Data Analysis**.

Additionally, Fastweb offers **FastwebAI Agents**, a workflow editor for AI agents, developed in collaboration with **Vellum**.

🔗 **Access FastwebAI Agents:** <https://fastwebai-agents.fastedge.it/>

📖 **Vellum Documentation:** <https://docs.vellum.ai/home/getting-started/overview>

📄 **Internal FAQ:** available in the dedicated Google Docs directory.



Usage Examples – Fastweb AI Factory Models

Below you will find practical examples (in **cURL** and **Python**) for the main model categories: **Embeddings, Chat, LLM, Audio, and Vision**.



1. jina-embeddings-v3

📖 Description:

Text **embedding model** optimized for semantic search and retrieval. Converts sentences into numerical vectors (*floats*) useful for:

- semantic search
- clustering
- classification



Example cURL

```
curl -X POST https://bpod1.ai-factory.fastweb.it/v1/embeddings \
-H "Authorization: Bearer $TOKEN" \
-H "Content-Type: application/json" \
-d '{
  "input": "Text to convert into embeddings",
  "model": "jinaai/jina-embeddings-v3",
  "encoding_format": "float"
}'
```

▶ Example Python

```
from openai import OpenAI
client = OpenAI(base_url="https://bpod1.ai-factory.fastweb.it/v1", api_key=TOKEN)

response = client.embeddings.create(
    model="jinaai/jina-embeddings-v3",
    input="Text to convert into embeddings",
    encoding_format="float"
)

print(response.data[0].embedding)
```

◆ 2. MIIA (Fastweb)

📖 Description:

A **conversational chat model** developed by Fastweb, optimized for natural and fluent responses in Italian.

▶ Example cURL

```
curl -X POST https://bpod1.ai-factory.fastweb.it/v1/chat/completions \
-H "Authorization: Bearer $TOKEN" \
-H "Content-Type: application/json" \
-d '{
    "model": "Fastweb/FastwebMIIA-7B",
    "messages": [{"role": "user", "content": "Hi, what can you do?"}],
    "max_tokens": 200
}'
```

▶ Example Python

```
from openai import OpenAI
client = OpenAI(base_url="https://bpod1.ai-factory.fastweb.it/v1", api_key=TOKEN)

response = client.chat.completions.create(
    model="Fastweb/FastwebMIIA-7B",
    messages=[{"role": "user", "content": "Hi, what can you do?"}],
    max_tokens=200
)

print(response.choices[0].message["content"])
```

◆ 3. LLaMA 3.3 70B (Meta AI)

📖 Description:

A powerful **open-source LLM** from Meta AI, suitable for advanced text generation and **function calling**.

▶ Example cURL

```
curl -X POST https://bpod1.ai-factory.fastweb.it/v1/chat/completions \
-H "Authorization: Bearer $TOKEN" \
-H "Content-Type: application/json" \
-d '{
  "model": "meta/llama-3.3-70b-instruct",
  "messages": [
    {"role": "user", "content": "What's the weather like in Rome?"}
  ],
  "tools": [
    {
      "type": "function",
      "function": {
        "name": "get_weather",
        "description": "Gets the current weather conditions for a city",
        "parameters": {
          "type": "object",
          "properties": {
            "city": {"type": "string", "description": "City name, e.g.
Milan"}
          },
          "required": ["city"]
        }
      }
    }
  ],
  "tool_choice": "auto",
  "max_tokens": 200
}'
```

▶ Example Python

```
from openai import OpenAI
client = OpenAI(base_url="https://bpod1.ai-factory.fastweb.it/v1", api_key=TOKEN)

response = client.chat.completions.create(
    model="meta/llama-3.3-70b-instruct",
    messages=[{"role": "user", "content": "Explain Einstein's theory of
relativity"}],
    max_tokens=300
)

print(response.choices[0].message["content"])
```

◆ 4. Whisper (OpenAI)

📖 Description:

A multilingual **speech-to-text** model, highly accurate even with background noise or strong accents.

▶ Example cURL

```
curl -X POST https://bpod1.ai-factory.fastweb.it/v1/audio/transcriptions \
-H "Authorization: Bearer $TOKEN" \
-H "Content-Type: multipart/form-data" \
-F "file=@audio.wav" \
-F "model=openai/whisper-large-v3-turbo"
```

▶ Example Python

```
from openai import OpenAI
client = OpenAI(base_url="https://bpod1.ai-factory.fastweb.it/v1", api_key=TOKEN)

with open("audio.wav", "rb") as f:
    response = client.audio.transcriptions.create(
        model="openai/whisper-large-v3-turbo",
        file=f
    )

print(response.text)
```

◆ 5. Gemma Vision (Google)

📖 Description:

A **multimodal model** that combines text and image analysis. Enables visual reasoning combined with text queries.

▶ Example cURL

```
curl -X POST https://bpod1.ai-factory.fastweb.it/v1/chat/completions \
-H "Authorization: Bearer $TOKEN" \
-H "Content-Type: application/json" \
-d '{
    "model": "google/gemma-3-27b-it",
    "messages": [
        {"role": "user", "content": [
            {"type": "text", "text": "What do you see in this image?"},
            {"type": "image_url", "image_url": {"url":
"https://example.com/image.jpg"}}
        ]}
    ]
}'
```

▶ Example Python

```
from openai import OpenAI
client = OpenAI(base_url="https://bpod1.ai-factory.fastweb.it/v1", api_key=TOKEN)

response = client.chat.completions.create(
    model="google/gemma-3-27b-it",
    messages=[
```

```
        {"role": "user", "content": [  
            {"type": "text", "text": "what do you see in this image?"},  
            {"type": "image_url", "image_url": {"url":  
"https://example.com/image.jpg"}}  
        ]  
    ]  
)  
  
print(response.choices[0].message["content"])
```